

# Shiny app deployment using ShinyProxy

Tobia De Koninck – Open Analytics  
September 22, 2023

# Open Analytics

# Data Science Company



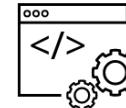
# Data Science Company



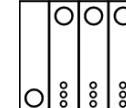
STATISTICAL  
CONSULTING



SCIENTIFIC  
PROGRAMMING



APPLICATION  
DEVELOPMENT AND  
INTEGRATION



DATA ANALYSIS  
HARDWARE AND  
HOSTING

# Agenda

Challenges when deploying Shiny apps in production:

- containers
- installing R packages and system dependencies
- using apps with different R, Bioconductor or Shiny versions
- persistence & databases
- authentication & authorisation
- monitoring
- initialization times

Goal: give you inspiration and tips on what is possible with Shiny and ShinyProxy

# ShinyProxy

- deploy data science apps in an enterprise context
- Shiny, Dash, Streamlit, Jupyter notebooks, IDEs, Gradio, Voila, Flask ...
- supports standard authentication/authorization technologies
- supports high availability
- scalable

<https://shinyproxy.io>



# Scalability & use-cases

- single (Docker) server with one app and a few users
- public instance (cluster) with 50 apps and 300 users per day
- cluster with 200 users providing R and Python IDEs
- validated (GxP) environment with full audit logging

# Containers

# Containers



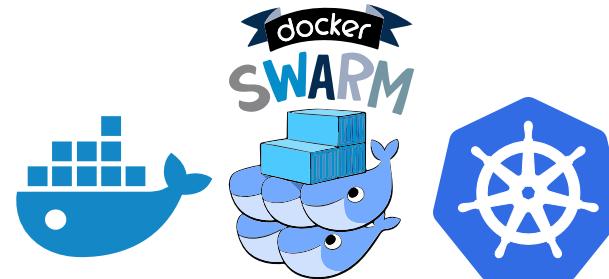
# Docker

Docker makes it easy to create, deploy, and run applications using containers. It allows to combine an application with everything it needs, such as libraries, dependencies, configuration, ... into a container image.

The **Dockerfile** is a set of instructions telling Docker how to create the Docker image. For example, which R version you need, which R packages etc.

# Why use Docker in ShinyProxy?

- the application is bundled and installed into a Docker image
- every time a user runs an application, a container spins up and serves the application
- fully isolated 'workspace' per session
- use different R, Shiny or Bioconductor versions on a single server, without installing this upfront on the server
- full control on memory and cpu usage
- archive every version of an application in order to achieve reproducibility
- Docker, Docker Swarm and Kubernetes supported



Extensible architecture: working on additional (cloud) container runtimes

# Experience with Docker

Who is already using Docker? What do you like?

# Experience with Docker

There is a learning curve, but in general we notice data scientists are able to quickly learn and use Docker in a good way.

In addition, because Docker makes it easier to reproduce results, it's very useful in (data) science.

Many data science departments are adopting containers.

# Installing R packages and system dependencies

# Working Example

In order to guide the discussion, we will use a working example:

interactiveDisplay

| Package for enabling powerful shiny web displays of Bioconductor objects

<https://www.bioconductor.org/packages/release/bioc/html/interactiveDisplay.html>

# Create the Dockerfile

```
FROM rocker/r-ver:4.3.1

RUN apt-get update && \
    apt-get install -y --no-install-recommends libz-dev liblzma-dev libbz2-dev \
    libcurl4-openssl-dev libxt6

RUN R -q -e 'install.packages("BiocManager"); BiocManager::install(version = "3.17")'
RUN R -q -e 'BiocManager::install(c("GenomicRanges", "Gviz"), version = "3.17")'
RUN R -q -e 'BiocManager::install(c("ggbio", "interactiveDisplay"), version = "3.17")'

COPY Rprofile.site /usr/local/lib/R/etc/
EXPOSE 3838
CMD ["R", "-q", "-e", "library(interactiveDisplay);data(mmgr);display(mmgr);"]
```

# System Dependencies

- find system dependencies of R packages on: <https://packagemanager.posit.co/client/#/repos/2/packages/shiny>
- overview of common system dependencies in R: <https://r-universe.dev/sysdeps/>
- all Ubuntu packages can be found on <https://packages.ubuntu.com/>
- <https://rocker-project.org/use/extending.html>

Often trial-and-error: both for CRAN and Bioconductor packages.

# r2u

Fast, easy and good?

Full integration with apt as every binary resolves all its dependencies: No more installations (of pre-built archives) only to discover that a shared library is missing. No more surprises.

<https://eddelbuettel.github.io/r2u/>

Very promising solution.

Supports Bioconductor

# r2u

```
FROM rocker/r2u:jammy

RUN R -q -e 'install.packages("BiocManager"); BiocManager::install(version = "3.17")'
RUN R -q -e 'BiocManager::install(c("GenomicRanges", "Gviz"), version = "3.17")'
RUN R -q -e 'BiocManager::install(c("ggbio", "interactiveDisplay"), version = "3.17")'

COPY Rprofile.site /usr/lib/R/etc/
EXPOSE 3838
CMD ["R", "-q", "-e", "library(interactiveDisplay);data(mmgr);display(mmgr);"]
```

Easy: ✓

Fast: ✓ 1m36s vs 34m27s

Good: ✓

# Pak

pak installs R packages from CRAN, Bioconductor, GitHub, URLs, git repositories, local files and directories. It is an alternative to `install.packages()` and `devtools::install_github()`. pak is fast, safe and convenient.

<https://pak.r-lib.org/>

- Supports Bioconductor
- installs system dependencies as well: <https://pak.r-lib.org/dev/reference/sysreqs.html>
- automatically selects the Bioconductor version using the R version.
- ⚠️ pak did not install all required system dependencies -> trial and error still needed
- automatically parallelize package download, installation and compilation
- (slower than r2u: 13m14s)

# Pak

```
FROM rocker/r-ver:4.3.1

RUN apt-get update && \
    apt-get install -y --no-install-recommends liblzma-dev libbz2-dev libxt6

RUN R -q -e 'install.packages("pak")'
RUN R -q -e 'pak::pkg_install(c("GenomicRanges", "Gviz", "ggbio", "interactiveDisplay"))'

COPY Rprofile.site /usr/local/lib/R/etc/
EXPOSE 3838
CMD ["R", "-q", "-e", "library(interactiveDisplay);data(mmgr);display(mmgr);"]
```

# Renv

The `renv` package helps you create reproducible environments for your R projects. Use `renv` to make your R projects more isolated, portable and reproducible.

Reproducible: `renv` records the exact package versions you depend on, and ensures those exact versions are the ones that get installed wherever you go.

<https://rstudio.github.io/renv/>

- Supports Bioconductor
- can be combined with `pak` (although it is still very new and experimental)
- possible to select the Bioconductor version
- Dockerfile can be very similar for completely different projects

<https://rstudio.github.io/renv/articles/bioconductor.html>

# Renv

```
FROM rocker/r-ver:4.3.1

RUN apt-get update && \
    apt-get install -y --no-install-recommends libbz-dev libbz2-dev liblzma-dev \
    libcurl4-openssl-dev libxt6

RUN R -q -e 'install.packages("renv")'
RUN R -q -e 'renv:::renv_pak_init(stream = "devel", force = TRUE)'
RUN mkdir /app

COPY renv.lock /app
WORKDIR /app

RUN R -q -e 'renv::restore()'

COPY Rprofile.site /usr/lib/R/etc/
EXPOSE 3838
CMD ["R", "-q", "-e", "library(interactiveDisplay);data(mmgr);display(mmgr);"]
```

# Installing packages and system dependencies

What tool do you use?

Which tool have you already tested?



# Build the Dockerfile

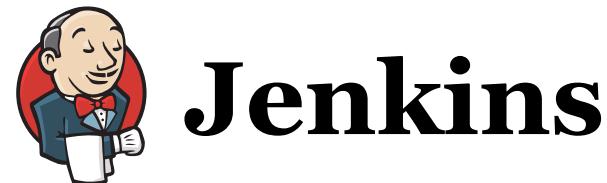
```
docker build -t myregistry.com/interactive-display  
docker push myregistry.com/interactive-display
```

Works on Linux, MacOS and Windows.

# Build the Dockerfile

Building and pushing can and should be fully automated:

- store the Dockerfile in Git
- edit the file
- git commit && git push
- automatically start building the image



# ShinyProxy configuration

# Update the ShinyProxy config

Edit the application.yml file and restart ShinyProxy.

As usual, best practice to automate this:

- store the application.yml in Git
- edit the file
- git commit && git push
- automatically deploy a new ShinyProxy server without downtime

specs:

- id: interactiveDisplay  
display-name: InteractiveDisplay  
description: package for enabling powerful shiny web displays of Bioconductor objects  
container-image: myregistry.com/interactive-display

# Apps with different R, Bioconductor or Shiny versions

In essence this is solved by using Docker images

Let's deploy another app using an older R version:

```
FROM openanalytics/r-ver:3.6.3

RUN apt-get update && apt-get install -y --no-install-recommends \
    libz-dev

RUN R -q -e "install.packages('shiny')"

CMD ["R", "-q", "-e", "shiny::runExample('01_hello')"]
```

# Reproducibility

- using Docker images already provides a lot of reproducibility: create once, run infinite times
- reproducing the Docker image itself: pin all versions in the Dockerfile
- allow users to run older versions of your app: GxP / clinical trials
- run multiple versions side-by-side
- related: when reproducibility is important, so are automatic tests

specs:

```
- id: reproducibility
  max-instances: -1
  container-image: |
    "myregistry.com/img:#{proxy.getRuntimeObject('SHINYPROXY_PARAMETERS').getValue('version')}"
  parameters:
    definitions:
      - id: version
        default-value: v1.0.1
  value-sets:
    - values:
        version:
          - v1.0.1
          - v1.0.0
          - v0.9.0
```

# Reproducibility

How do you guarantee reproducibility when using Shiny apps?  
Do you write tests for your Shiny app?

# Persistence & databases

# Persistence for Applications

Many Shiny apps require a directory to store files.

ShinyProxy allows accessing a shared directory using volumes.

- shared between multiple users:

```
specs:  
  - id: interactiveDisplay  
    container-image: myregistry.com/interactive-display  
    container-volumes: ["/mnt/interactive-display:/mnt/interactive-display"]
```

- unique to every user:

```
specs:  
  - id: interactiveDisplay  
    container-image: myregistry.com/interactive-display  
    container-volumes: ["/mnt/interactive-display/#{proxy.userId}:/mnt/interactive-display"]
```

- essential when using ShinyProxy for hosting IDEs (Rstudio, Jupyter notebook ...)

# Persistence for Applications

- use cloud-based filesystems, like AWS EFS or Azure FileShare
- supports PV & PVCs on Kubernetes
- often lacks file-locking support, e.g. SQLite -> use a network database
- file system permission can be tricky
- same mechanisms can be used for ephemeral storage

Do you face any other challenges with persistence storage?

# Interaction with databases

- Shiny app can directly access any database service, using regular R code
- because Docker isolates the network, accessing databases on the same server as ShinyProxy can be tricky, but is possible
- best practice to pass database host, username and password as environment variables:

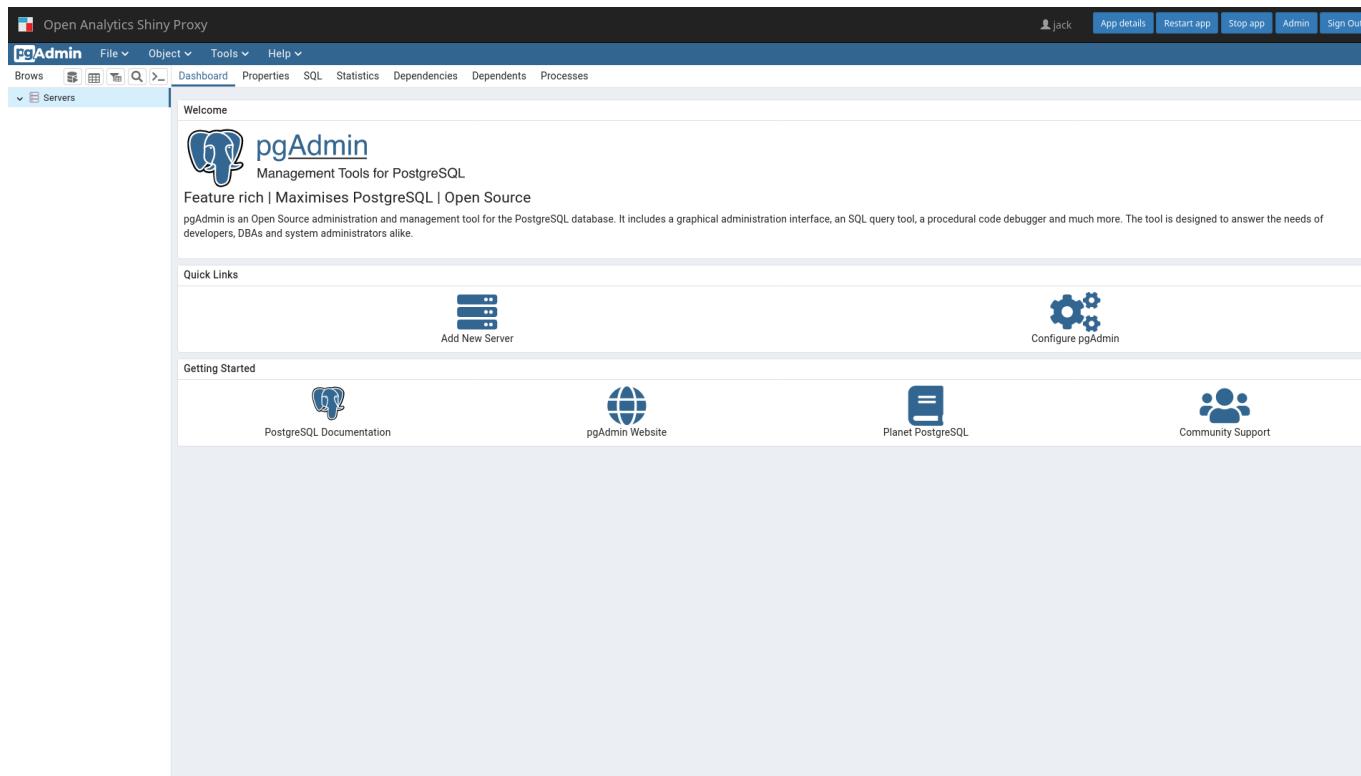
```
specs:  
  - id: interactiveDisplay  
    container-image: myregistry.com/interactive-display  
    container-env:  
      DB_HOST: "172.16.10.10"  
      DB_USERNAME: "postgres"  
      DB_PASSWORD: "myPassword!"
```

- Kubernetes secrets natively supported
- Docker swarm secrets natively supported
- ⚠ do not store the password in the R code or on git

# Interaction with databases

Bonus: run pgadmin on ShinyProxy:

<https://github.com/openanalytics/shinyproxy-pgadmin-demo>

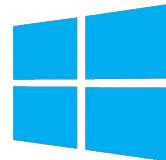


# Authentication & Authorisation

# Authentication

## Proving a user's identity

- best-practice: authentication should not happen in the Shiny app
- don't reinvent the wheel: OpenID Connect, LDAP, SAML
- integrating with external tools allows for e-mail verification, resetting passwords, 2FA ...
- ShinyProxy authenticates the user and forwards this information to the container using environment variables



Active Directory

# Authentication

```
library(shiny)

shinyServer(function(input, output, session) {

  output$username <- renderText({ Sys.getenv("SHINYPROXY_USERNAME") })
  output$groups <- renderText({ Sys.getenv("SHINYPROXY_USERGROUPS") })

})
```

# Authorisation

Giving a user permission to access or do something

ShinyProxy: whether a user can access an app:

```
specs:  
- id: auth-demo  
  display-name: Demo of authentication and authorisation  
  container-image: myregistry.com/auth-demo  
  access-users:  
    - jack  
  access-groups:  
    - mathematicians
```

# Authorisation

Shiny: whether a user can see some data or perform a certain action:

# Monitoring

# Logs & Metrics

<https://shinyproxy.io/documentation/usage-statistics/>

<https://github.com/openanalytics/shinyproxy-monitoring>

Usage statistics, logging and audit logging supported.



# Initialization time

# Problem

- a new container is created when the user opens an app
- the container is created for a specific user -> can be customized for this user
- user has to wait for container and the R process to start

# Solution

Solution: container pre-initialization

- ShinyProxy maintains a pool of containers, which are ready to use
- automatically scale-up and scale-down this pool
- automatically re-create containers when the configuration changes
- will support multiple users using a single container

# Thanks

- <https://shinyproxy.io/>
- <https://shinyproxy.io/documentation/demos/>
- Feature requests or bug reports: <https://github.com/openanalytics/shinyproxy>
- Support: <https://support.openanalytics.eu/>
- <https://openanalytics.eu/>
- [info@openanalytics.eu](mailto:info@openanalytics.eu)